

# Neural Networks

Matt R

February 23, 2022

## Contents

<b>1</b>	<b>Logistic Regression</b>	<b>2</b>
1.1	The Gradient . . . . .	3
1.1.1	Vectorization in Python . . . . .	4
<b>2</b>	<b>Neural Networks: A Single Hidden Layer</b>	<b>9</b>

# 1 Logistic Regression

We begin with a review of binary classification and logistic regression. To this end, suppose we have training examples  $x \in \mathbb{R}^{m \times n}$  with binary labels  $y \in \{0, 1\}^{1 \times n}$ . We desire to train a model which yields an output  $a$  which represents

$$a = \mathbb{P}(y = 1|x).$$

To this end, let  $\sigma : \mathbb{R} \rightarrow (0, 1)$  denote the sigmoid function, i.e.,

$$\sigma(z) = \frac{1}{1 + e^{-z}},$$

and let  $w \in \mathbb{R}^m$ ,  $b \in \mathbb{R}$ , and let

$$a = \sigma(w^T x + b).$$

To analyze the accuracy of model, we need a way to compare  $y$  and  $a$ , and ideally this functional comparison can be optimized with respect to  $(w, b)$  in such a way to minimize the error. To this end, we note that

$$\mathbb{P}(y|x) = a^y(1 - a)^{1-y},$$

or rather

$$\mathbb{P}(y = 1|x) = a, \quad \mathbb{P}(y = 0|x) = 1 - a,$$

so  $\mathbb{P}(y|x)$  represents the corrected probability. Now since we want

$$a \approx 1 \quad \text{when } y = 1,$$

and

$$a \approx 0 \quad \text{when } y = 0,$$

and  $0 \leq a \leq 1$ , any error using differences won't be refined enough to analyze when tuning the model. Moreover, since introducing the sigmoid function, our usual mean-squared-error function won't be convex. This leads us to apply the log function, which when restricted to  $(0, 1)$  is a bijective mapping of  $(0, 1) \rightarrow (-\infty, 0)$ . This leads us to define our log-loss function

$$\begin{aligned} \mathbb{L}(a, y) &= -\log(\mathbb{P}(y|x)) \\ &= -\log(a^y(1 - a)^{1-y}) \\ &= -[y \log(a) + (1 - y) \log(1 - a)], \end{aligned}$$

and finally, since we wish to analyze how our model performs on the entire training set, we need to average our log-loss functions to obtain our cost function  $\mathbb{J}$  defined by

$$\begin{aligned}\mathbb{J}(w, b) &= \frac{1}{n} \sum_{j=1}^n \mathbb{L}(a_j, y_j) \\ &= -\frac{1}{n} \sum_{j=1}^n [y_j \log(a_j) + (1 - y_j) \log(1 - a_j)] \\ &= -\frac{1}{n} \sum_{j=1}^n [y_j \log(\sigma(w^T x_j + b)) + (1 - y_j) \log(1 - \sigma(w^T x_j + b))] .\end{aligned}$$

## 1.1 The Gradient

To compute the gradient of our cost function  $\mathbb{J}$ , we first write  $\mathbb{J}$  as a sum of compositions as follows: We have the log-loss function considered as a map  $\mathbb{L} : (0, 1) \times \mathbb{R} \rightarrow \mathbb{R}$ ,

$$\mathbb{L}(a, y) = -[y \log(a) + (1 - y) \log(1 - a)] ,$$

we have the sigmoid function  $\sigma : \mathbb{R} \rightarrow (0, 1)$  with  $\sigma(z) = a$  and  $\sigma'(z) = a(1 - a)$ , and we have the collection of affine-functionals  $\phi_x : \mathbb{R}^m \times \mathbb{R} \rightarrow \mathbb{R}$  given by

$$\phi_x(w, b) = w^T x + b,$$

for which we fix an arbitrary  $x \in \mathbb{R}^m$  and write  $\phi = \phi_x$ , and set  $z = \phi(w, b)$ . Finally, we introduce the auxiliary function  $\mathcal{L} : \mathbb{R}^m \times \mathbb{R} \rightarrow \mathbb{R}$  given by

$$\mathcal{L}(w, b) = \mathbb{L}(\sigma(\phi(w, b)), y).$$

Then by the chain rule, we have that

$$\begin{aligned}d\mathcal{L} &= d_a \mathbb{L}(a, y) \circ d\sigma(z) \circ d_w \phi(w, b) \\ &= \left[ -\frac{y}{a} + \frac{1-y}{1-a} \right] \cdot a(1-a) \cdot [x^T \quad 1] \\ &= [-y(1-a) + a(1-y)] \cdot [x^T \quad 1] \\ &= (a-y) [x^T \quad 1]\end{aligned}$$

Composition turns into matrix multiplication in the tangent space.

Moreover, since in Euclidean space, we have that  $\nabla f = (df)^T$ , and hence that

$$\nabla \mathcal{L}(w, b) = (a - y) \begin{bmatrix} x \\ 1 \end{bmatrix},$$

or rather

$$\partial_w \mathbb{L}(a, y) = (a - y)x, \quad \partial_b \mathbb{L}(a, y) = a - y.$$

Finally, since our cost function  $\mathbb{J}$  is the sum-log-loss, we have by linearity that

$$\begin{aligned} \partial_w \mathbb{J}(w, b) &= \frac{1}{n} \sum_{j=1}^n (a_j - y_j) x_j \\ &= \frac{1}{n} ((a - y) \cdot x^T)^T \\ &= \frac{1}{n} x \cdot (a - y)^T \end{aligned}$$

and

$$\partial_b \mathbb{J}(w, b) = \frac{1}{n} \sum_{j=1}^n (a_j - y_j).$$

### 1.1.1 Vectorization in Python

Here we include the general code to train a model using logistic regression without regularization and without tuning on a cross-validation set.

```

1 import copy
2
3 import numpy as np
4
5 def sigmoid(z):
6     """
7     Parameters
8     -----
9     z : array_like
10
11     Returns
12     -----
13     sigma : array_like
14     """
15
16     sigma = (1 / (1 + np.exp(-z)))
17     return sigma
18

```

```

19 def cost_function(x, y, w, b):
20     """
21     Parameters
22     -----
23     x : array_like
24         x.shape = (m, n) with m-features and n-examples
25     y : array_like
26         y.shape = (1, n)
27     w : array_like
28         w.shape = (m, 1)
29     b : float
30
31     Returns
32     -----
33     J : float
34         The value of the cost function evaluated at (w, b)
35     dw : array_like
36         dw.shape = w.shape = (m, 1)
37         The gradient of J with respect to w
38     db : float
39         The partial derivative of J with respect to b
40     """
41
42     # Auxiliary assignments
43     m, n = x.shape
44     z = w.T @ x + b
45     assert z.size == n
46     a = sigmoid(z).reshape(1, n)
47     dz = a - y
48
49     # Compute cost J
50     J = (-1 / n) * (np.log(a) @ y.T + np.log(1 - a) @ (1 - y).T)
51
52     # Compute dw and db
53     dw = (x @ dz.T) / m
54     assert dw.shape == w.shape
55     db = np.sum(dz) / m
56
57     return J, dw, db
58
59 def grad_descent(x, y, w, b, alpha=0.001, num_iters=2000, print_cost=False):
60     """
61     Parameters
62     -----
63     x, y, w, b : See cost_function above for specifics.
64         w and b are chosen to initialize the descent (likely all components 0)
65     alpha : float

```

```

66         The learning rate of gradient descent
67     num_iters : int
68         The number of times we wish to perform gradient descent
69
70     Returns
71     -----
72     costs : List[float]
73         For each iteration we record the cost-values associated to (w, b)
74     params : Dict[w : array_like, b : float]
75         w : array_like
76             Optimized weight parameter w after iterating through grad descent
77         b : float
78             Optimized bias parameter b after iterating through grad descent
79     grads : Dict[dw : array_like, db : float]
80         dw : array_like
81             The optimized gradient with respect to w
82         db : float
83             The optimized derivative with respect to b
84     """
85
86     costs = []
87     w = copy.deepcopy(w)
88     b = copy.deepcopy(b)
89     for i in range(num_iters):
90         J, dw, db = cost_function(x, y, w, b)
91         w = w - alpha * dw
92         b = b - alpha * db
93
94         if i % 100 == 0:
95             costs.append(J)
96             if print_cost:
97                 idx = int(i / 100) - 1
98                 print(f'Cost_after_iteration_{i}:_{costs[idx]}')
99
100     params = {'w' : w, 'b' : b}
101     grads = {'dw' : dw, 'db' : db}
102
103     return costs, params, grads
104
105 def predict(w, b, x):
106     """
107     Parameters
108     -----
109     w : array_like
110         w.shape = (m, 1)
111     b : float
112     x : array_like

```

```

113         x.shape = (m, n)
114
115     Returns
116     -----
117     y_predict : array_like
118         y_pred.shape = (1, n)
119         An array containing the prediction of our model applied to training
120         data x, i.e., y_pred = 1 or y_pred = 0.
121     """
122
123     m, n = x.shape
124     # Get probability array
125     a = sigmoid(w.T @ x + b)
126     # Get boolean array with False given by a < 0.5
127     pseudo_predict = ~(a < 0.5)
128     # Convert to binary to get predictions
129     y_predict = pseudo_predict.astype(int)
130
131     return y_predict
132
133 def model(x_train, y_train, x_test, y_test, alpha=0.001, num_iters=2000, accuracy=True)
134     """
135     Parameters:
136     -----
137     x_train, y_train, x_test, y_test : array_like
138         x_train.shape = (m, n_train)
139         y_train.shape = (1, n_train)
140         x_test.shape = (m, n_test)
141         y_test.shape = (1, n_test)
142     alpha : float
143         The learning rate for gradient descent
144     num_iters : int
145         The number of times we wish to perform gradient descent
146     accuracy : Boolean
147         Use True to print the accuracy of the model
148
149     Returns:
150     d : Dict
151         d['costs'] : array_like
152             The costs evaluated every 100 iterations
153         d['y_train_preds'] : array_like
154             Predicted values on the training set
155         d['y_test_preds'] : array_like
156             Predicted values on the test set
157         d['w'] : array_like
158             Optimized parameter w
159         d['b'] : float

```

```

160         Optimized parameter b
161         d['learning_rate'] : float
162         The learning rate alpha
163         d['num_iters'] : int
164         The number of iterations with which gradient descent was performed
165
166     """
167
168     m = x_train.shape[0]
169     # initialize parameters
170     w = np.zeros((m, 1))
171     b = 0.0
172     # optimize parameters
173     costs, params, grads = grad_descent(x_train, y_train, w, b, alpha, num_iters)
174     w = params['w']
175     b = params['b']
176     # record predictions
177     y_train_preds = predict(w, b, x_train)
178     y_test_preds = predict(w, b, x_test)
179     # group results into dictionary for return
180     d = {'costs' : costs,
181         'y_train_preds' : y_train_preds,
182         'y_test_preds' : y_test_preds,
183         'w' : w,
184         'b' : b,
185         'learning_rate' : alpha,
186         'num_iters' : num_iters}
187
188     if accuracy:
189         train_acc = 100 - np.mean(np.abs(y_train_preds - y_train)) * 100
190         test_acc = 100 - np.mean(np.abs(y_test_preds - y_test)) * 100
191         print(f'Training_Accuracy:_{train_acc}%')
192         print(f'Test_Accuracy:_{test_acc}%')
193
194
195     return d

```



## 2 Neural Networks: A Single Hidden Layer

Suppose we wish to consider the binary classification problem given the training set  $(x, y)$  with  $x \in \mathbb{R}^{s_0 \times n}$  and  $y \in \{0, 1\}^n$ . Usually with logistic regression we have the following type of structure:

$$[x^1, \dots, x^{s_0}] \xrightarrow{\varphi} [z] \xrightarrow{g} [a] \xrightarrow{=} \hat{y},$$

where

$$z = \varphi(x) = w^T x + b,$$

is our affine-linear transformation, and

$$a = g(z) = \sigma(z)$$

is our sigmoid function. Such a structure will be called a *network*, and the  $[a]$  is known as the *activation node*. Logistic regression can be too simplistic of a model for many situations. To modify this model to handle more complex situations, we introduce a new “hidden layer” of nodes with their own (possibly different) activation functions. That is, we consider a network of the following form:

$$\underbrace{\begin{bmatrix} x^1 \\ \vdots \\ x^{s_0} \end{bmatrix}}_{\text{Layer 0}} \xrightarrow{\varphi^{[1]}} \underbrace{\begin{bmatrix} z^{[1]1} \\ \vdots \\ z^{[1]s_1} \end{bmatrix}}_{\text{Layer 1}} \xrightarrow{g^{[1]}} \underbrace{\begin{bmatrix} a^{[1]1} \\ \vdots \\ a^{[1]s_1} \end{bmatrix}}_{\text{Layer 2}} \xrightarrow{\varphi^{[2]}} \underbrace{\begin{bmatrix} z^{[2]} \\ \vdots \\ z^{[2]} \end{bmatrix}}_{\text{Layer 2}} \xrightarrow{g^{[2]}} \underbrace{\begin{bmatrix} a^{[2]} \\ \vdots \\ a^{[2]} \end{bmatrix}}_{\text{Layer 2}} \xrightarrow{=} \hat{y},$$

where

$$\begin{aligned} \varphi^{[1]} : \mathbb{R}^{s_0} &\rightarrow \mathbb{R}^{s_1}, & \varphi^{[1]}(x) &= W^{[1]}x + b^{[1]}, \\ \varphi^{[2]} : \mathbb{R}^{s_1} &\rightarrow \mathbb{R}, & \varphi^{[2]}(x) &= W^{[2]}x + b^{[2]}, \end{aligned}$$

and  $W^{[1]} \in \mathbb{R}^{s_1 \times s_0}$ ,  $W^{[2]} \in \mathbb{R}^{1 \times s_1}$ ,  $b^{[1]} \in \mathbb{R}^{s_1}$ ,  $b^{[2]} \in \mathbb{R}$ , and  $g^{[\ell]}$  is a *broadcasted* activator function (e.g., the sigmoid function  $\sigma(z)$ , or  $\tanh(z)$ , or  $\text{ReLU}(z)$ ). Such a network is called a 2-layer neural network where  $x$  is the input layer (called layer-0),  $a^{[1]}$  is a hidden layer (called layer-1), and  $a^{[2]}$  is the output layer (called layer-2).

**Definition 2.1.** Suppose  $g : \mathbb{R} \rightarrow \mathbb{R}$  is any function. Then we say  $\bar{g} : \mathbb{R}^{m \times n} \rightarrow \mathbb{R}^{m \times n}$  is the **broadcast** of  $g$  if

$$\begin{aligned} \bar{g}(A) &= \bar{g}(A_j^i e_i^j) \\ &= g(A_j^i) e_i^j, \end{aligned}$$

where  $A \in \mathbb{R}^{m \times n}$  and  $\{e_i^j : 1 \leq i \leq m, 1 \leq j \leq n\}$  is the standard basis for  $\mathbb{R}^{m \times n}$ .

Let us lay out all of these functions explicitly (in the Smooth Category) as to facilitate our later computations for our cost function and our gradients. To this end:

$$\begin{aligned} \varphi^{[1]} : \mathbb{R}^{s_0} &\rightarrow \mathbb{R}^{s_1}, & d\varphi^{[1]} : T\mathbb{R}^{s_0} &\rightarrow T\mathbb{R}^{s_1}, \\ z^{[1]} = \varphi^{[1]}(x) &= W^{[1]}x + b^{[1]}, & d\varphi_x^{[1]}(v) &= W^{[1]}v; \end{aligned}$$

$$\begin{aligned} g^{[1]} : \mathbb{R}^{s_1} &\rightarrow \mathbb{R}^{s_1}, & dg^{[1]} : T\mathbb{R}^{s_1} &\rightarrow T\mathbb{R}^{s_1}, \\ a^{[1]} = g^{[1]}(z^{[1]}), & dg_{z^{[1]}}^{[1]}(v) = \begin{bmatrix} g^{[1]'}(z^{[1]1}) & 0 & \cdots & 0 \\ 0 & g^{[1]'}(z^{[1]2}) & \cdots & 0 \\ \vdots & \ddots & \ddots & \vdots \\ 0 & 0 & \cdots & g^{[1]'}(z^{[1]s_1}) \end{bmatrix} \begin{bmatrix} v^1 \\ v^2 \\ \vdots \\ v^{s_1} \end{bmatrix}; \end{aligned}$$

$$\begin{aligned} \varphi^{[2]} : \mathbb{R}^{s_1} &\rightarrow \mathbb{R}, & d\varphi^{[2]} : T\mathbb{R}^{s_1} &\rightarrow T\mathbb{R}, \\ z^{[2]} = \varphi^{[2]}(a^{[1]}) &= W^{[2]}a^{[1]} + b^{[2]}, & d\varphi_{a^{[1]}}^{[2]}(v) &= W^{[2]}v; \end{aligned}$$

$$\begin{aligned} g^{[2]} : \mathbb{R} &\rightarrow \mathbb{R}, & dg^{[2]} : T\mathbb{R} &\rightarrow T\mathbb{R}, \\ a^{[2]} = g^{[2]}(z^{[2]}), & dg_{z^{[2]}}^{[2]}(v) = g^{[2]'}(z^{[2]}) \cdot v. \end{aligned}$$

That is, given an input  $x \in \mathbb{R}^{s_0}$ , we get a predicted value  $\hat{y}$  of the form

$$\hat{y} = g^{[2]} \circ \varphi^{[2]} \circ g^{[1]} \circ \varphi^{[1]}(x).$$

This compositional function is known as *forward propagation*.

Since we wish to optimize our model with respect to our parameter  $W^{[\ell]}$  and  $b^{[\ell]}$ , we consider a generic loss function  $\mathbb{L} : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ ,  $\mathbb{L}(\hat{y}, y)$ , and by acknowledging the potential abuse of notation, we assume  $y$  is fixed, and consider the aforementioned as a function of a single-variable

$$\mathbb{L}_y : \mathbb{R} \rightarrow \mathbb{R}, \quad \mathbb{L}_y(\hat{y}) = \mathbb{L}(\hat{y}, y).$$

We now define the compositional function

$$F : \mathbb{R}^{s_0} \rightarrow \mathbb{R}, \quad F(x) = \mathbb{L}_y \circ g^{[2]} \circ \varphi^{[2]} \circ g^{[1]} \circ \varphi^{[1]}(x).$$

As we mentioned before, we wish to optimize with respect to our parameters, but our above composition doesn't make this dependence explicit for computations. To this end, we first previously considered the generic affine-linear transformation

$$\varphi : \mathbb{R}^m \rightarrow \mathbb{R}^k, \quad \varphi(x) = Wx + b,$$

with  $W \in \mathbb{R}^{k \times m}, b \in \mathbb{R}^k$ . We now change our point-of-view and consider the related

$$\phi : \mathbb{R}^{k \times m} \times \mathbb{R}^k \rightarrow \mathbb{R}^k, \quad \phi(W, b) = Wx + b,$$

for some fixed  $x \in \mathbb{R}^m$ . Then we see that

$$d\phi : T\mathbb{R}^{k \times m} \times T\mathbb{R}^k \rightarrow T\mathbb{R}^k,$$

$$\begin{aligned} d\phi_{(W,b)}(V, v) &= \left. \frac{d}{dt} \right|_{t=0} \phi(W + tV, b + tv) \\ &= \left. \frac{d}{dt} \right|_{t=0} (W + tV)x + (b + tv) \\ &= Vx + v. \end{aligned}$$

This leads to the further decomposition of the form

$$\phi(W, b) = \psi(W) + \mathbb{1}_{\mathbb{R}^k}(b),$$

where  $\mathbb{1}_{\mathbb{R}^k} : \mathbb{R}^k \rightarrow \mathbb{R}^k$  is the identity function, and  $\psi : \mathbb{R}^{k \times m} \rightarrow \mathbb{R}^k$  is given by

$$\psi(W) = Wx.$$

Then by the above computation, we have that

$$d\psi_W(V) = Vx.$$

Moreover, suppose

$$\{E_1^1, E_1^2, \dots, E_1^m, E_2^1, E_2^2, \dots, E_2^m, \dots, E_k^1, E_k^2, \dots, E_k^m\}$$

is an ordering of the standard basis  $\{E_\alpha^\beta\}$  for  $\mathbb{R}^{k \times m}$  with

$$[E_\alpha^\beta]_j^i = \delta_\alpha^i \delta_j^\beta,$$

and

$$V = V_j^i E_i^j,$$

then  $d\psi_W \in \mathbb{R}^{k \times (m+k)}$  with matrix-representation

$$d\psi_W = \begin{bmatrix} x^1 & x^2 & \cdots & x^m & 0 & 0 & \cdots & 0 & \cdots & 0 \\ 0 & \cdots & 0 & 0 & & & & & & \\ 0 & 0 & \cdots & 0 & x^1 & x^2 & \cdots & x^m & \cdots & 0 \\ 0 & \cdots & 0 & 0 & & & & & & \end{bmatrix}$$

Taking this further, we now consider the map

$$\Phi : \mathbb{R}^{k \times m} \times \mathbb{R}^k \times \mathbb{R}^m \rightarrow \mathbb{R}^k, \quad \Phi(W, b, x) = Wx + b,$$

and then computing our differential for

$$d\Phi : T\mathbb{R}^{k \times m} \times T\mathbb{R}^k \times T\mathbb{R}^m \rightarrow T\mathbb{R}^k,$$

yields

$$\begin{aligned} d\Phi_{(W,b,x)}(V, v, p) &= \left. \frac{d}{dt} \right|_{t=0} \Phi(W + tV, b + tv, x + tp) \\ &= \left. \frac{d}{dt} \right|_{t=0} ((W + tV)(x + tp) + (b + tv)) \\ &= \left. \frac{d}{dt} \right|_{t=0} (Wx + tVx + tWv + t^2Vp + b + tv) \\ &= Vx + v + Wp \\ &= d\phi_{(W,b)}(V, v) + d\varphi_x(p) \end{aligned}$$

This function  $\Phi$  is what we want in our compositional-function, and so we redefine  $F$  as

$$F(W^{[2]}, b^{[2]}, W^{[1]}, b^{[1]}, x) = \mathbb{L}_y \circ g^{[2]} \circ \Phi^{[2]} \circ (W^{[2]}, b^{[2]}, g^{[1]} \circ \Phi^{[1]} \circ (W^{[1]}, b^{[1]}, x))$$

Taking the exterior derivative, and noting the composition turns into matrix multiplication on the tangent space, we get

$$\begin{aligned} dF_{(W^{[2]}, b^{[2]}, W^{[1]}, b^{[1]}, x)}(U, u, V, v, p) &= d(\mathbb{L}_y)_{a^{[2]}} \cdot dg_{z^{[2]}}^{[2]} \cdot d\Phi_{(W^{[2]}, b^{[2]}, a^{[1]})}^{[2]} \cdot (U, u, dg_{z^{[1]}}^{[1]} \cdot d\Phi_{(W^{[1]}, b^{[1]}, x)}^{[1]}(V, v, p)) \\ &= d(\mathbb{L}_y)_{a^{[2]}} \cdot dg_{z^{[2]}}^{[2]} \cdot (d\phi_{(W^{[2]}, b^{[2]})}^{[2]}(U, u) + d\varphi_{a^{[1]}}^{[2]} \cdot dg_{z^{[1]}}^{[1]} \cdot (d\phi_{(W^{[1]}, b^{[1]})}^{[1]}(V, v) + d\varphi_x^{[1]}(p))) \\ &= d(\mathbb{L}_y)_{a^{[2]}} \cdot dg_{z^{[2]}}^{[2]} \cdot d\phi_{(W^{[2]}, b^{[2]}, \{a^{[1]}\})}^{[2]}(U, u) \\ &\quad + d(\mathbb{L}_y)_{a^{[2]}} \cdot dg_{z^{[2]}}^{[2]} \cdot d\varphi_{a^{[1]}, \{W^{[2]}, b^{[2]}\}}^{[2]} \cdot dg_{z^{[1]}}^{[1]} \cdot d\phi_{(W^{[1]}, b^{[1]}, \{x\})}^{[1]}(V, v) \\ &\quad + d(\mathbb{L}_y)_{a^{[2]}} \cdot dg_{z^{[2]}}^{[2]} \cdot d\varphi_{a^{[1]}, \{W^{[2]}, b^{[2]}\}}^{[2]} \cdot dg_{z^{[1]}}^{[1]} \cdot d\varphi_{x, \{W^{[1]}, b^{[1]}\}}^{[1]}(p) \\ &=: dF^{[2]} + dF^{[1]} + dF^{[0]}, \end{aligned}$$

where  $dF^{[2]}$  represents the differential with respect to the parameters going from layer-1 to layer-2,  $dF^{[1]}$  represents the differential with respect to the parameters going from layer-0 to layer-1, and  $dF^{[0]}$  represents the differential with respect to  $x$ .

Recalling that the gradient is the transpose of the exterior derivative in Euclidean space, we then conclude that

$$\begin{aligned}\nabla F &= (dF)^T \\ &= \left( dF^{[2]} + dF^{[1]} + dF^{[0]} \right)^T \\ &= \nabla F^{[2]} + \nabla F^{[1]} + \nabla F^{[0]},\end{aligned}$$

and respectively,

$$\nabla F^{[2]} = \left( d(\mathbb{L}_y)_{a^{[2]}} \cdot dg_{z^{[2]}}^{[2]} \cdot d\phi_{(W^{[2]}, b^{[2]}, \{a^{[1]}\})}^{[2]} \right)^T$$